

On-device AI coding. No cloud. No limits.

A complete coding agent that executes entirely on your machine. No API calls. No usage caps.

[Join Waitlist](#)[Our Approach](#)

ZERO TELEMETRY

NATIVE INFERENCE

100% OFFLINE

YOUR HARDWARE, YOUR RULES

NO TOKENS, NO LIMITS

SPECIALIZED SLM

UNBOUNDED CONTEXT

THE PROBLEM

You don't own your AI. And you're being watched.



DATA EXTRACTION

001

They train on your code.

Every prompt. Every file. Every fix.
It flows through infrastructure you don't control –
improving systems they want to use to replace you.

ARTIFICIAL SCARCITY

002

They meter your ambition.

Slowdowns, overages, caps.
Right when you're deep in a sprint, the meter decides
you've had enough.

SILENT DOWNGRADES

003

They change the model.

They silently downgrade to cheaper models during peak
load. Full price, degraded experience.

CLOUD DEPENDENCY

004

They control your flow.

Every completion makes a round trip across the internet.
Thousands of tiny interruptions, every single day.

INTRODUCING RIG

Everything local. Own your AI.

A complete AI coding agent running entirely on your own hardware. No usage limits. No cloud dependency.





OFFLINE

Work offline

Flights. Spotty Wi-Fi. Network outages. Nothing stops your flow.

UNLIMITED

Remove the meter

Refactor the whole codebase. Riff on an idea all day. Run agent loops without thinking about cost.

PRIVACY

Sever the connection

Your code, keystrokes, and files never leave your machine. Not anonymized. Not aggregated. Not sent.

LATENCY

Stop waiting

No round-trip to a data center. Inference happens on your machine, in single-digit milliseconds.

OUR APPROACH

Purpose beats scale.

Rig is a closed system — model, context, tools, and inference — engineered together for one job: real coding work.

STEP 01

A focused model, trained specifically for coding.

STEP 02


Full intelligence, compressed to fit your machine.

STEP 03

A custom runtime, engineered for Apple Silicon.

The model runs through a custom inference engine optimized specifically for Apple Silicon. Model, context engine, and tools are designed as a single coordinated system.

That tight integration is what makes local execution fast, reliable, and practical.



Metric	Rig	Cloud APIs
First token latency	300 ms	400-1,000 ms
Cost per 1K tokens	\$0.00	\$0.01-0.06

CAPABILITIES

Your machine, unleashed.

[01]

Understands your architecture.

Builds a connected model of modules, dependencies, and relationships so reasoning happens across files and aligns with your architecture.

[02]

Tracks relationships, prevents breakage.

Edits that respect function contracts, type boundaries, and dependency graphs — reducing bugs and regressions.

[03]

Strategizes before acting.

Explore → Plan → Execute workflows ensure multiple steps are reasoned out before changes occur.

[04]

Executes complex coding workflows.

From refactors to test generation to feature builds — coordinate tools, code edits, web search, and commands as needed.

[05]

Isolates agent sandboxes.

Each agent runs in its own workspace so experiments are safe, parallel workflows don't clash, and code changes stay isolated until you merge them.

[06]

Runs at full speed.

Custom Rust inference engine optimized for CUDA and Metal — delivering up to 144 tokens per second on consumer hardware.

LATENCY 0ms No round-trip required	PRIVACY 100% Air-gapped by design	COST / TOKEN \$0 Your GPU, your tokens	UPTIME Local No dependency on cloud
---	--	---	--

III ENGINEERED INTELLIGENCE

Built for control freaks

CUSTOM MODEL
OPTIMIZED FOR CONSUMER
HARDWARE

INFERENCE
CROSS-OS USING RUST

CONTEXT GRAPH
REPO-WIDE CODE
UNDERSTANDING

RIG

> Scanning hardware...

> Found WA · 16GB RAM

> Loading RIG Model OK

> Indexing 2,418 files · 87,102 symbols

✓ Ready. Network: OFF · Telemetry: OFF

λ refac

TERMINAL UI
BUILT IN RUST AND
BLAZING FAST

HEAVILY TUNED
CONSISTENT TOOL CALLS
AND PLAN USE

OPINIONATED
FOCUSED ON CODE
CORRECTNESS

NEURAL ENGINE · RIG-B00 · LOCAL GPU

🔓 EARLY ACCESS

Rig is almost ready.

We're inviting engineers to run it on
real code and help shape what ships.

you@company.com

JOIN WAITLIST

🔗 FAQ

Frequently asked questions.

- 01

What is Rig?

▼
- 02

What Model does Rig use?

▼
- 03

What are the Hardware Requirements?

▼
- 04

How Does Rig Compare to Large Cloud Models?

▼
- 05

Can Rig Search the Web?

▼
- 06

How will Rig be Priced?

▼

07 Will Rig Collect my Data?



08 When Will Rig be Available?



Break free from big AI

Request Early Access ↗

No credit card. No usage meter.

RIG

Local-first AI for developers who refuse to be
dependent on the cloud.

CONNECT

[Twitter](#)
[LinkedIn](#)

LEGAL

[Privacy Policy](#)

© 2020 RIG AI INC. ALL RIGHTS RESERVED.

■ ALL SYSTEMS LOCAL